



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Rust programming language in high-performance computing

Rust Meetup

Michał Sudwoj

06.05.2022


Rust programming language in high-performance computing

2022-05-06

Rust programming language in
high-performance computing

Rust Meetup
Michał Sudwoj
06.05.2022

Before we begin...



2022-05-06

Rust programming language in high-performance computing

- └ Before we begin...
- └ Before we begin...



About me



- currently:
 - MSc Student in CSE at ETH Zürich
 - Data Engineer at Bank Vontobel
- before:
 - Intern at CSCS
 - Data Scientist at zkipster
 - BSc in CSE at ETH Zürich

2022-05-06

Rust programming language in high-performance computing

└ About me

└ About me

1. CSE: Computational Science and Engineering
2. from quantum chemistry to astrophysics
3. biology, meteorology, geology, ...

About me



- currently:
 - * MSc Student in CSE at ETH Zürich
 - * Data Engineer at Bank Vontobel
- before:
 - * Intern at CSCS
 - * Data Scientist at zkipster
 - * BSc in CSE at ETH Zürich

What about you?

2022-05-06

Rust programming language in high-performance computing

└ About me

What about you?

1. Software Engineers?
2. R&D?
3. PDE? Finite difference?

CSCS Internship: Part 1

In the first part, the Rust will be evaluated for potential usage at CSCS. In particular, the following questions should be tackled:

- how to **install and run Rust programs** with "user access" rights on Piz Daint
- is MPI wrapper for Rust compatible with Cray's implementation
- how to **interface Rust program with numerical libraries**, such as MKL, MAGMA, ScaLAPACK, cuBlasXt, etc.
- how to **write GPU-enabled application in Rust**; what are the complications or simplifications comparing to the C/C++/FORTRAN GPU applications
- how to debug and profile Rust-based programs
- how rich is the functionality of Rust, e.g. the availability of special mathematical functions, support of matrices or multi-dimensional arrays, etc.

2022-05-06

Rust programming language in high-performance computing

└ CSCS Internship

└ Part 1

└ CSCS Internship: Part 1

CSCS Internship: Part 1

In the first part, the Rust will be evaluated for potential usage at CSCS. In particular, the following questions should be tackled:

- how to **install and run Rust programs** with "user access" rights on Piz Daint
- is MPI wrapper for Rust compatible with Cray's implementation
- how to **interface Rust program with numerical libraries**, such as MKL, MAGMA, ScaLAPACK, cuBlasXt, etc.
- how to **write GPU-enabled application in Rust**; what are the complications or simplifications comparing to the C/C++/FORTRAN GPU applications
- how to debug and profile Rust-based programs
- how rich is the functionality of Rust, e.g. the availability of special mathematical functions, support of matrices or multi-dimensional arrays, etc.

CSCS Internship: Part 2

In the second part, a **performance comparison between Rust and C/C++/FORTRAN** will be conducted, by idiomatically implementing a parallel distributed linear algebra algorithm or a scientific mini-app code in the target languages. The performance analysis is **not only limited to computational performance, but may include analysis of other factors, such as ease of implementation, number of bugs made, testability, readability, maintainability**, etc.

2022-05-06

Rust programming language in high-performance computing

└ CSCS Internship

└ Part 2

└ CSCS Internship: Part 2

CSCS Internship: Part 2

In the second part, a **performance comparison between Rust and C/C++/FORTRAN** will be conducted, by idiomatically implementing a parallel distributed linear algebra algorithm or a scientific mini-app code in the target languages. The performance analysis is **not only limited to computational performance, but may include analysis of other factors, such as ease of implementation, number of bugs made, testability, readability, maintainability**, etc.

Part 1: Installation

```
> curl https://sh.rustup.rs -sSf | sh
> rustup toolchain install nightly
> rustup target add nvptx64-nvidia-cuda
> # On Piz Daint
> export
> CARGO_TARGET_X86_64_UNKNOWN_LINUX_GNU_RUSTFLAGS="
>     -C target-cpu=haswell
>     -C relocation-model=dynamic-no-pic
> "
> CARGO_TARGET_NVPTX64_NVIDIA_CUDA_RUSTFLAGS="
>     -C target-cpu=sm_60
>     -C target-feature=+sm_60,+ptx60
>     -C relocation-model=dynamic-no-pic
> "
> MPICC=cc
> cargo install ptx-linker
```

Rust programming language in high-performance computing

Part 1

Installation

Part 1: Installation

2022-05-06

Part 1: Installation

```
> curl https://sh.rustup.rs -sSf | sh
> rustup toolchain install nightly
> rustup target add nvptx64-nvidia-cuda
> # On Piz Daint
> export
> CARGO_TARGET_X86_64_UNKNOWN_LINUX_GNU_RUSTFLAGS="
>     -C target-cpu=haswell
>     -C relocation-model=dynamic-no-pic
> "
> CARGO_TARGET_NVPTX64_NVIDIA_CUDA_RUSTFLAGS="
>     -C target-cpu=sm_60
>     -C target-feature=+sm_60,+ptx60
>     -C relocation-model=dynamic-no-pic
> "
> MPICC=cc
> cargo install ptx-linker
```

Part 1: MPI

mpi crate: <https://github.com/rsmpi/rsmpi>

- thin wrapper around MPI implementation
- works with OpenMPI and MPICH
- works out of the box
- **derive macro for custom types**

```
#[derive(Equivalence)]
```

```
struct Particle {
```

```
    mass: f64,
```

```
    charge: i8,
```

```
    kind: Kind,
```

```
}
```

Rust programming language in high-performance computing

Part 1

MPI

Part 1: MPI

2022-05-06

Part 1: MPI

mpi crate: <https://github.com/rsmpi/rsmpi>

- thin wrapper around MPI implementation
 - works with OpenMPI and MPICH
 - works out of the box
 - **derive macro for custom types**
- ```
#[derive(Equivalence)]
struct Particle {
 mass: f64,
 charge: i8,
 kind: Kind,
}
```



## Part 1: Interfacing with C

- bindgen crate
- existing bindings to BLAS, LAPACK, cublas, MPI, ...

2022-05-06

Rust programming language in high-performance computing

└─ Part 1

└─ Interfacing with C

└─ Part 1: Interfacing with C

Part 1: Interfacing with C

- bindgen crate
- existing bindings to BLAS, LAPACK, cublas, MPI, ...

## Part 2: Toy Problem

- fourth-order numerical diffusion in  $xy$ -plane

$$\frac{\partial \phi}{\partial t} = -\alpha_4 \nabla_{xy}^4 \phi = \underbrace{-\alpha_4 \Delta_{xy}^2 \phi}_{\text{inline}} = \underbrace{-\alpha_4 \Delta_{xy} (\Delta_{xy} \phi)}_{\text{laplap}}$$

- on unit cube
- boundary conditions:  $\partial_{\Omega} = 0$

## Rust programming language in high-performance computing

2022-05-06

└ Part 2

└ Toy Problem

└ Part 2: Toy Problem

Part 2: Toy Problem

- fourth-order numerical diffusion in  $xy$ -plane

$$\frac{\partial \phi}{\partial t} = -\alpha_4 \nabla_{xy}^4 \phi = \underbrace{-\alpha_4 \Delta_{xy}^2 \phi}_{\text{inline}} = \underbrace{-\alpha_4 \Delta_{xy} (\Delta_{xy} \phi)}_{\text{laplap}}$$

- on unit cube
- boundary conditions:  $\partial_{\Omega} = 0$

- Xue (2000)
- Used in weather simulations as a smoothing kernel
- Reference implementation available - Oliver Fuhrer

$$\phi_{i,j,k}^{n+1} = \phi_{i,j,k}^n - \frac{\alpha_4 \Delta t}{\Delta x \Delta y} \Delta \left( -4\phi_{i,j,k}^n + \phi_{i-1,j,k}^n + \phi_{i+1,j,k}^n + \phi_{i,j-1,k}^n + \phi_{i,j+1,k}^n \right)$$

2022-05-06


Rust programming language in high-performance computing

└ Part 2

└ laplap

$$\phi_{i,j,k}^{n+1} = \phi_{i,j,k}^n - \frac{\alpha_4 \Delta t}{\Delta x \Delta y} \Delta \left( -4\phi_{i,j,k}^n + \phi_{i-1,j,k}^n + \phi_{i+1,j,k}^n + \phi_{i,j-1,k}^n + \phi_{i,j+1,k}^n \right)$$

# laplap




2022-05-06

Rust programming language in high-performance computing

└ Part 2  
└ laplap  
└ laplap

laplap



inline

$$\begin{aligned}\phi_{i,j,k}^{n+1} = & \phi_{i,j,k}^n - \frac{\alpha_4 \Delta t}{(\Delta x)^2 (\Delta y)^2} ( \\ & - 20\phi_{i,j,k}^n \\ & + 8\phi_{i-1,j,k}^n + 8\phi_{i+1,j,k}^n + 8\phi_{i,j-1,k}^n + 8\phi_{i,j+1,k}^n \\ & - 2\phi_{i-1,j-1,k}^n - 2\phi_{i-1,j+1,k}^n - 2\phi_{i+1,j-1,k}^n - 2\phi_{i+1,j+1,k}^n \\ & - \phi_{i-2,j,k}^n - \phi_{i+2,j,k}^n - \phi_{i,j-2,k}^n - \phi_{i,j+2,k}^n )\end{aligned}$$

Rust programming language in high-performance computing


2022-05-06

└ Part 2  
└ laplap  
└ inline

```
inline
phi_{i,j,k}^{n+1} = phi_{i,j,k}^n - \frac{\alpha_4 \Delta t}{(\Delta x)^2 (\Delta y)^2} (
 - 20\phi_{i,j,k}^n
 + 8\phi_{i-1,j,k}^n + 8\phi_{i+1,j,k}^n + 8\phi_{i,j-1,k}^n + 8\phi_{i,j+1,k}^n
 - 2\phi_{i-1,j-1,k}^n - 2\phi_{i-1,j+1,k}^n - 2\phi_{i+1,j-1,k}^n - 2\phi_{i+1,j+1,k}^n
 - \phi_{i-2,j,k}^n - \phi_{i+2,j,k}^n - \phi_{i,j-2,k}^n - \phi_{i,j+2,k}^n)
```

1. What do you think - which version is faster?

inline




Rust programming language in high-performance computing

2022-05-06


Part 2

laplap

inline




# Initial conditions




2022-05-06

- Rust programming language in high-performance computing
  - Part 2
    - Initial conditions
      - Initial conditions




# Result



2022-05-06

Rust programming language in high-performance computing

- Part 2
  - Result
    - Result





# Measurement setup

- Fortran, C++, Rust
- Algorithms: laplap, inline
- Toolchains
  - Fortran, C++: GNU, Cray, Intel, PGI
  - Rust: rustc
- Sequential
- Parallel
  - Fortran, C++: OpenMP
  - Rust: Rayon
- GPU
  - Fortran, C++: OpenACC, OpenMP offloading, CUDA
  - Rust: Accel

## Rust programming language in high-performance computing

2022-05-06

### Part 2

#### Measurement setup

#### Measurement setup

Measurement setup

- Fortran, C++, Rust
- Algorithms: laplap, inline
- Toolchains
  - Fortran, C++: GNU, Cray, Intel, PGI
  - Rust: rustc
- Sequential
- Parallel
  - Fortran, C++: OpenMP
  - Rust: Rayon
- GPU
  - Fortran, C++: OpenACC, OpenMP offloading, CUDA
  - Rust: Accel

1. MPI works, but not enough time to debug the partitioner

## Measurement setup

⇒ 84 language-compiler-algorithm combinations

32 versions did not compile or run :(

⇒ 52 versions × 4 grid sizes

⇒ 524 measurements (1, 2, 4, 8, 12 CPU cores; GPU)

2022-05-06

Rust programming language in high-performance computing

Part 2

Measurement setup

Measurement setup

Measurement setup

⇒ 84 language-compiler-algorithm combinations

32 versions did not compile or run :(

⇒ 52 versions × 4 grid sizes

⇒ 524 measurements (1, 2, 4, 8, 12 CPU cores; GPU)

- arrays in column-major order
- -O3 or equivalent
- no LTO
- optimization reports
- shared libraries
- C interface

```
void diffuse(
 float * in_field,
 float * out_field,
 size_t nx,
 size_t ny,
 size_t nz,
 size_t num_halo,
 float alpha,
 size_t num_iter
)
```

## Rust programming language in high-performance computing

2022-05-06

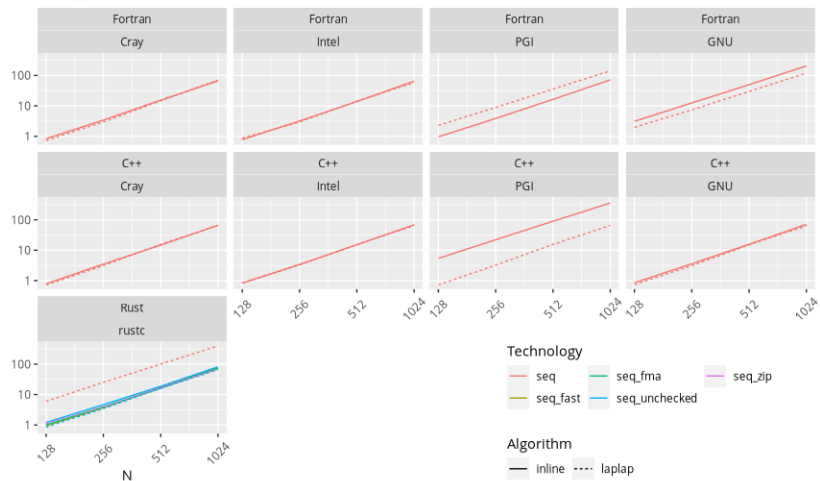
### Part 2

#### Measurement setup

```
■ arrays in column-major order
■ -O3 or equivalent
■ no LTO
■ optimization reports
■ shared libraries
■ C interface
void diffuse(
 float * in_field,
 float * out_field,
 size_t nx,
 size_t ny,
 size_t nz,
 size_t num_halo,
 float alpha,
 size_t num_iter
)
```

#### 1. Show code on gitlab


Runtime (s)




2022-05-06

# Rust programming language in high-performance computing

## Results



Runtime (s)




Technology — seq — seq\_fast — seq\_fma — seq\_unchecked — seq\_zip

Algorithm — inline — laplap

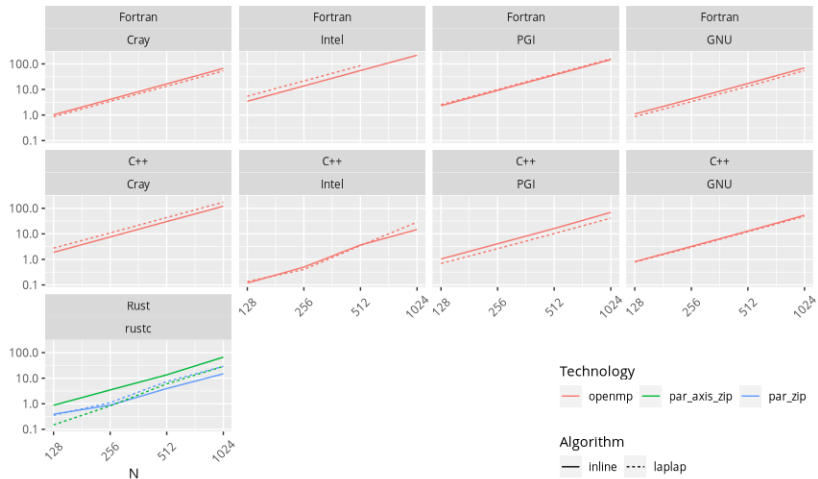
# Rust programming language in high-performance computing

## Results

2022-05-06




Runtime (s)




2022-05-06

# Rust programming language in high-performance computing

## Results



Runtime (s)




Technology — accel — cuda — openmp\_target

Algorithm — inline - - - laplap

# Rust programming language in high-performance computing

## Results

2022-05-06



# Conclusions

- + Rust fast enough for scientific software
- + Safer code
- + Clearer error messages
- Support for certain features is lacking
- ~ All languages sometimes lead to lots of boilerplate

2022-05-06

Rust programming language in high-performance computing

└─ Conclusions

└─ Conclusions

1. Rust compiler does good job of vectorizing code
2. Rust: clear what you get
3. Fortran/C++: different support, many bugs
4. GNU doesn't warn about lack of offloading
5. Cray sometimes silently generates invalid PTX code

Conclusions

- + Rust fast enough for scientific software
- + Safer code
- + Clearer error messages
- Support for certain features is lacking
- ~ All languages sometimes lead to lots of boilerplate



## Recommendations

- in 2020
  - Rust for frontend/driver
- hopes for after 2020
  - GPU support
  - Evolve ndarray
  - ScaLAPACK bindings, ...
- $\Rightarrow$  continue to pursue Rust in HPC
  - it has potential

See:

- <https://www.arewelearningyet.com/scientific-computing/>
- <https://www.arewelearningyet.com/gpu-computing/>

2022-05-06

Rust programming language in high-performance computing

└─ Conclusions

└─ Recommendations

└─ Recommendations

Recommendations

- in 2020
  - Rust for frontend/driver
- hopes for after 2020
  - GPU support
  - Evolve ndarray
  - ScaLAPACK bindings, ...
- $\Rightarrow$  continue to pursue Rust in HPC
  - it has potential

See:

- <https://www.arewelearningyet.com/scientific-computing/>
- <https://www.arewelearningyet.com/gpu-computing/>

1. With Cray and Intel moving to LLVM, cross-language LTO soon?

Two years later ...

2022-05-06

Rust programming language in high-performance computing  
└ Two years later ...

Two years later ...

not much has changed :(

2022-05-06

Rust programming language in high-performance computing  
└ Two years later ...

not much has changed :(

## Two years later ...

- cargo-cmake → corrosion
- rust-cuda → rust-gpu
- nvptx64-nvidia-cuda still a Tier 2 target

2022-05-06

Rust programming language in high-performance computing


└ Two years later ...

└ Two years later ...

Two years later ...

- cargo-cmake → corrosion
- rust-cuda → rust-gpu
- nvptx64-nvidia-cuda still a Tier 2 target

Questions?



2022-05-06

Rust programming language in high-performance computing

└─ Questions?

└─ Questions?

Questions?



## Keep in touch?

- `michal@sudwoj.name`
- `https://github.com/westernmagic/rust-in-hpc`
- `www.linkedin.com/in/michalsudwoj`

2022-05-06

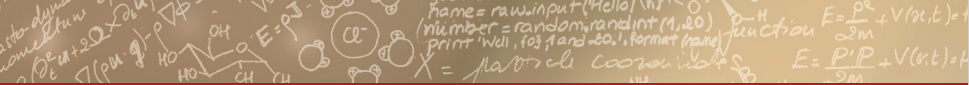
Rust programming language in high-performance computing

└ Keep in touch?

└ Keep in touch?

Keep in touch?

- `michal@sudwoj.name`
- `https://github.com/westernmagic/rust-in-hpc`
- `www.linkedin.com/in/michalsudwoj`



Thank you for you attention

2022-05-06

Rust programming language in high-performance computing  
Keep in touch?

Thank you for you attention